

# wrangle\_act

June 28, 2022

## 1 Project: Wrangling and Analyze Data

### 1.1 Data Gathering

In the cell below, gather **all** three pieces of data for this project and load them in the notebook. **Note:** the methods required to gather each data are different. 1. Directly download the WeRate-Dogs Twitter archive data (twitter\_archive\_enhanced.csv)

```
In [1]: import pandas as pd
import requests
import os
import tweepy
import json
import re
import matplotlib.pyplot as plt
%matplotlib inline
```

```
In [111]: #Reading the csv file into a dataframe
archive_df = pd.read_csv("twitter-archive-enhanced.csv")
```

2. Use the Requests library to download the tweet image prediction (image\_predictions.tsv)

```
In [3]: url = "https://d17h27t6h515a5.cloudfront.net/topher/2017/August/599fd2ad_image-predictions/image_predictions.tsv"
response = requests.get(url)
folder_name = 'images'
# Make directory if it doesn't already exist
if not os.path.exists(folder_name):
    os.makedirs(folder_name)
# Make a path to the file and write to it
with open(os.path.join(folder_name, url.split('/')[-1]), mode = 'wb') as file:
    file.write(response.content)
```

```
In [4]: # Reading the tsv file into a dataframe
image_df = pd.read_csv("images/image_predictions.tsv", sep='\t')
```

3. Use the Tweepy library to query additional data via the Twitter API (tweet\_json.txt)

```

In [6]: consumer_key = ''
        consumer_secret = ''
        access_token = ''
        access_secret = ''

        auth = tweepy.OAuthHandler(consumer_key, consumer_secret)
        auth.set_access_token(access_token, access_secret)

        api = tweepy.API(auth, wait_on_rate_limit=True, wait_on_rate_limit_notify=False, compress=True)

In [7]: id = archive_df['tweet_id']
        tweet_list = []
        id_errors = []

        for number in id:
            try:
                tweet_text = api.get_status(number, tweet_mode='extended')._json
                tweet_list.append({'tweet_id': tweet_text['id'],
                                  'favorite_count': tweet_text['favorite_count'],
                                  'retweet_count': tweet_text['retweet_count']})
            except Exception as e:
                id_errors.append(e)

In [8]: # writing the queried tweets to a text file.
        with open('tweet_json.txt', 'w') as file:
            file.write(json.dumps(tweet_list))

In [9]: # Creating a dataframe of columns tweet id, favorite count and retweet count.
        tweet_df = pd.DataFrame(tweet_list, columns=['tweet_id', 'favorite_count', 'retweet_count'])

```

## 1.2 Assessing Data

In this section, detect and document at least **eight (8) quality issues** and **two (2) tidiness issue**. You must use **both** visual assessment and programmatic assessment to assess the data.

**Note:** pay attention to the following key points when you access the data.

- You only want original ratings (no retweets) that have images. Though there are 5000+ tweets in the dataset, not all are dog ratings and some are retweets.
- Assessing and cleaning the entire dataset completely would require a lot of time, and is not necessary to practice and demonstrate your skills in data wrangling. Therefore, the requirements of this project are only to assess and clean at least 8 quality issues and at least 2 tidiness issues in this dataset.
- The fact that the rating numerators are greater than the denominators does not need to be cleaned. This [unique rating system](#) is a big part of the popularity of WeRateDogs.
- You do not need to gather the tweets beyond August 1st, 2017. You can, but note that you won't be able to gather the image predictions for these tweets since you don't have access to the algorithm used.

## Visual assessment

```
In [112]: archive_df.head()
```

```
Out[112]:
```

	tweet_id	in_reply_to_status_id	in_reply_to_user_id	\
0	892420643555336193	NaN	NaN	
1	892177421306343426	NaN	NaN	
2	891815181378084864	NaN	NaN	
3	891689557279858688	NaN	NaN	
4	891327558926688256	NaN	NaN	

	timestamp	\
0	2017-08-01 16:23:56 +0000	
1	2017-08-01 00:17:27 +0000	
2	2017-07-31 00:18:03 +0000	
3	2017-07-30 15:58:51 +0000	
4	2017-07-29 16:00:24 +0000	

	source	\
0	<a href="http://twitter.com/download/iphone" r...	
1	<a href="http://twitter.com/download/iphone" r...	
2	<a href="http://twitter.com/download/iphone" r...	
3	<a href="http://twitter.com/download/iphone" r...	
4	<a href="http://twitter.com/download/iphone" r...	

	text	retweeted_status_id	\
0	This is Phineas. He's a mystical boy. Only eve...	NaN	
1	This is Tilly. She's just checking pup on you...	NaN	
2	This is Archie. He is a rare Norwegian Pouncin...	NaN	
3	This is Darla. She commenced a snooze mid meal...	NaN	
4	This is Franklin. He would like you to stop ca...	NaN	

	retweeted_status_user_id	retweeted_status_timestamp	\
0	NaN	NaN	
1	NaN	NaN	
2	NaN	NaN	
3	NaN	NaN	
4	NaN	NaN	

	expanded_urls	rating_numerator	\
0	https://twitter.com/dog_rates/status/892420643...	13	
1	https://twitter.com/dog_rates/status/892177421...	13	
2	https://twitter.com/dog_rates/status/891815181...	12	
3	https://twitter.com/dog_rates/status/891689557...	13	
4	https://twitter.com/dog_rates/status/891327558...	12	

	rating_denominator	name	doggo	floofer	pupper	puppo
0	10	Phineas	None	None	None	None
1	10	Tilly	None	None	None	None

2	10	Archie	None	None	None	None
3	10	Darla	None	None	None	None
4	10	Franklin	None	None	None	None

In [55]: image\_df.head()

```
Out[55]:
```

	tweet_id	jpg_url	\
0	666020888022790149	https://pbs.twimg.com/media/CT4udnOWwAA0aMy.jpg	
1	666029285002620928	https://pbs.twimg.com/media/CT42GRgUYAA5iDo.jpg	
2	666033412701032449	https://pbs.twimg.com/media/CT4521TWwAEvMyu.jpg	
3	666044226329800704	https://pbs.twimg.com/media/CT5Dr8HUEAA-lEu.jpg	
4	666049248165822465	https://pbs.twimg.com/media/CT5IQmsXIAAKY4A.jpg	

	img_num	p1	p1_conf	p1_dog	p2	\
0	1	Welsh_springer_spaniel	0.465074	True	collie	
1	1	redbone	0.506826	True	miniature_pinscher	
2	1	German_shepherd	0.596461	True	malinois	
3	1	Rhodesian_ridgeback	0.408143	True	redbone	
4	1	miniature_pinscher	0.560311	True	Rottweiler	

	p2_conf	p2_dog	p3	p3_conf	p3_dog
0	0.156665	True	Shetland_sheepdog	0.061428	True
1	0.074192	True	Rhodesian_ridgeback	0.072010	True
2	0.138584	True	bloodhound	0.116197	True
3	0.360687	True	miniature_pinscher	0.222752	True
4	0.243682	True	Doberman	0.154629	True

In [11]: tweet\_df.head()

```
Out[11]:
```

	tweet_id	favorite_count	retweet_count
0	892420643555336193	33798	7009
1	892177421306343426	29332	5302
2	891815181378084864	22056	3481
3	891689557279858688	36938	7227
4	891327558926688256	35312	7762

## Programmatic asesment

In [56]: archive\_df.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 2356 entries, 0 to 2355
Data columns (total 17 columns):
tweet_id                2356 non-null int64
in_reply_to_status_id   78 non-null float64
in_reply_to_user_id     78 non-null float64
timestamp               2356 non-null object
source                  2356 non-null object
text                    2356 non-null object
```

```

retweeted_status_id      181 non-null float64
retweeted_status_user_id 181 non-null float64
retweeted_status_timestamp 181 non-null object
expanded_urls            2297 non-null object
rating_numerator          2356 non-null int64
rating_denominator        2356 non-null int64
name                     2356 non-null object
doggo                    2356 non-null object
floofer                  2356 non-null object
pupper                   2356 non-null object
puppo                    2356 non-null object
dtypes: float64(4), int64(3), object(10)
memory usage: 313.0+ KB

```

```
In [57]: archive_df.describe()
```

```

Out[57]:
      tweet_id  in_reply_to_status_id  in_reply_to_user_id \
count  2.356000e+03                7.800000e+01          7.800000e+01
mean   7.427716e+17                7.455079e+17          2.014171e+16
std    6.856705e+16                7.582492e+16          1.252797e+17
min    6.660209e+17                6.658147e+17          1.185634e+07
25%    6.783989e+17                6.757419e+17          3.086374e+08
50%    7.196279e+17                7.038708e+17          4.196984e+09
75%    7.993373e+17                8.257804e+17          4.196984e+09
max    8.924206e+17                8.862664e+17          8.405479e+17

      retweeted_status_id  retweeted_status_user_id  rating_numerator \
count      1.810000e+02                1.810000e+02          2356.000000
mean      7.720400e+17                1.241698e+16          13.126486
std       6.236928e+16                9.599254e+16          45.876648
min       6.661041e+17                7.832140e+05           0.000000
25%       7.186315e+17                4.196984e+09          10.000000
50%       7.804657e+17                4.196984e+09          11.000000
75%       8.203146e+17                4.196984e+09          12.000000
max       8.874740e+17                7.874618e+17          1776.000000

      rating_denominator
count      2356.000000
mean       10.455433
std        6.745237
min         0.000000
25%        10.000000
50%        10.000000
75%        10.000000
max        170.000000

```

```
In [16]: sum(archive_df.duplicated())
```

```
Out[16]: 0
```

```
In [17]: archive_df.isnull().sum()
```

```
Out[17]: tweet_id          0
         in_reply_to_status_id    2278
         in_reply_to_user_id    2278
         timestamp              0
         source                 0
         text                   0
         retweeted_status_id    2175
         retweeted_status_user_id  2175
         retweeted_status_timestamp  2175
         expanded_urls          59
         rating_numerator       0
         rating_denominator     0
         name                   0
         doggo                  0
         floofer                0
         pupper                 0
         puppo                  0
         dtype: int64
```

```
In [58]: archive_df.name.value_counts().head()
```

```
Out[58]: None          745
         a              55
         Charlie        12
         Cooper         11
         Oliver         11
         Name: name, dtype: int64
```

```
In [19]: image_df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 2075 entries, 0 to 2074
Data columns (total 12 columns):
tweet_id    2075 non-null int64
jpg_url     2075 non-null object
img_num     2075 non-null int64
p1          2075 non-null object
p1_conf     2075 non-null float64
p1_dog      2075 non-null bool
p2          2075 non-null object
p2_conf     2075 non-null float64
p2_dog      2075 non-null bool
p3          2075 non-null object
p3_conf     2075 non-null float64
p3_dog      2075 non-null bool
```

```
dtypes: bool(3), float64(3), int64(2), object(4)
memory usage: 152.1+ KB
```

```
In [20]: image_df.describe()
```

```
Out[20]:
```

	tweet_id	img_num	p1_conf	p2_conf	p3_conf
count	2.075000e+03	2075.000000	2075.000000	2.075000e+03	2.075000e+03
mean	7.384514e+17	1.203855	0.594548	1.345886e-01	6.032417e-02
std	6.785203e+16	0.561875	0.271174	1.006657e-01	5.090593e-02
min	6.660209e+17	1.000000	0.044333	1.011300e-08	1.740170e-10
25%	6.764835e+17	1.000000	0.364412	5.388625e-02	1.622240e-02
50%	7.119988e+17	1.000000	0.588230	1.181810e-01	4.944380e-02
75%	7.932034e+17	1.000000	0.843855	1.955655e-01	9.180755e-02
max	8.924206e+17	4.000000	1.000000	4.880140e-01	2.734190e-01

```
In [21]: sum(image_df.duplicated())
```

```
Out[21]: 0
```

```
In [26]: tweet_df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 2327 entries, 0 to 2326
Data columns (total 3 columns):
tweet_id      2327 non-null int64
favorite_count 2327 non-null int64
retweet_count 2327 non-null int64
dtypes: int64(3)
memory usage: 54.6 KB
```

```
In [27]: sum(tweet_df.duplicated())
```

```
Out[27]: 0
```

## 1.2.1 Quality issues

### archive\_df

- There are columns with null values; 'in\_reply\_to\_status\_id', 'in\_reply\_to\_user\_id', 'retweeted\_status\_id', 'retweeted\_status\_user\_id', 'retweeted\_status\_timestamp', 'expanded\_urls'.
- The timestamp column should be of datetime type.
- There are plenty of dog names categorized as 'None', and others with 'a' and 'an'.
- There should be one column for all the dog stages.
- The tweet\_id column should be of object datatype.

## image\_df

- The text in the p1, p2 and p3 columns are a mix of upper and lower case characters.
- The tweet\_id column should be of object datatype.

#### tweet\_df

- The tweet\_id column should be of object datatype.

### 1.2.2 Tidiness

- The source column looks messy and should be dropped.
- The links in the text column should be removed.
- The three dataframes have similar properties and should be merged.

## 1.3 Cleaning Data

In this section, clean **all** of the issues you documented while assessing.

**Note:** Make a copy of the original data before cleaning. Cleaning includes merging individual pieces of data according to the rules of [tidy data](#). The result should be a high-quality and tidy master pandas DataFrame (or DataFrames, if appropriate).

```
In [113]: # Make copies of original pieces of data
          archive = archive_df.copy()
          images = image_df.copy()
          tweets = tweet_df.copy()
```

### 1.3.1 Issue #1:

**Define:** Drop columns with null values from the archive dataframe.

**Code:**

```
In [114]: archive.dropna(axis='columns', inplace=True)
```

```
In [139]: archive.drop(columns='source', inplace=True)
```

**Test**

```
In [140]: archive.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 2356 entries, 0 to 2355
Data columns (total 7 columns):
tweet_id          2356 non-null object
timestamp         2356 non-null datetime64[ns]
text              2356 non-null object
rating_numerator  2356 non-null int64
```



```
rating_denominator    2356 non-null int64
name                  2356 non-null object
dog-stage             2356 non-null object
dtypes: datetime64[ns](1), int64(2), object(4)
memory usage: 128.9+ KB
```

### 1.3.2 Issue #2:

**Define:** Convert the timestamp column to type datetime.

#### Code

```
In [116]: archive['timestamp'] = pd.to_datetime(archive['timestamp'])
```

#### Test

```
In [141]: archive.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 2356 entries, 0 to 2355
Data columns (total 7 columns):
tweet_id          2356 non-null object
timestamp         2356 non-null datetime64[ns]
text              2356 non-null object
rating_numerator  2356 non-null int64
rating_denominator 2356 non-null int64
name              2356 non-null object
dog-stage         2356 non-null object
dtypes: datetime64[ns](1), int64(2), object(4)
memory usage: 128.9+ KB
```

### 1.3.3 Issue #3

**Define:** Change the tweet\_id in the three dataframes to object datatype.

#### Code

```
In [118]: archive.tweet_id = archive.tweet_id.astype('str')
          images.tweet_id = images.tweet_id.astype('str')
          tweets.tweet_id = tweets.tweet_id.astype('str')
```

#### Test

```
In [142]: archive.info()
          images.info()
          tweets.info()
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 2356 entries, 0 to 2355
Data columns (total 7 columns):
tweet_id          2356 non-null object
timestamp         2356 non-null datetime64[ns]
text              2356 non-null object
rating_numerator  2356 non-null int64
rating_denominator 2356 non-null int64
name              2356 non-null object
dog-stage         2356 non-null object
dtypes: datetime64[ns](1), int64(2), object(4)
memory usage: 128.9+ KB
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 2075 entries, 0 to 2074
Data columns (total 12 columns):
tweet_id    2075 non-null object
jpg_url     2075 non-null object
img_num     2075 non-null int64
p1          2075 non-null object
p1_conf     2075 non-null float64
p1_dog      2075 non-null bool
p2          2075 non-null object
p2_conf     2075 non-null float64
p2_dog      2075 non-null bool
p3          2075 non-null object
p3_conf     2075 non-null float64
p3_dog      2075 non-null bool
dtypes: bool(3), float64(3), int64(1), object(5)
memory usage: 152.1+ KB
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 2327 entries, 0 to 2326
Data columns (total 3 columns):
tweet_id          2327 non-null object
favorite_count    2327 non-null int64
retweet_count     2327 non-null int64
dtypes: int64(2), object(1)
memory usage: 54.6+ KB

```

#### 1.3.4 Issue #4

**Define:** The dog names categorized as 'a' and 'an' should be added to the 'None' category instead.

#### Code

```
In [120]: archive['name'][archive.name.str.match('[a-z]+')] = 'None'
```

```
/opt/conda/lib/python3.6/site-packages/ipykernel_launcher.py:1: SettingWithCopyWarning:  
A value is trying to be set on a copy of a slice from a DataFrame
```

See the caveats in the documentation: <http://pandas.pydata.org/pandas-docs/stable/indexing.html#>

```
"""Entry point for launching an IPython kernel.
```

## Test

```
In [121]: archive.name.value_counts()
```

```
Out[121]: None      854  
         Charlie      12  
         Oliver      11  
         Lucy        11  
         Cooper      11  
         Tucker      10  
         Penny       10  
         Lola        10  
         Bo          9  
         Winston     9  
         Sadie       8  
         Daisy       7  
         Toby        7  
         Buddy       7  
         Bailey      7  
         Stanley     6  
         Koda        6  
         Leo         6  
         Jack        6  
         Rusty       6  
         Milo        6  
         Oscar       6  
         Jax         6  
         Dave        6  
         Bella       6  
         Scout       6  
         Bentley     5  
         George      5  
         Oakley      5  
         Chester     5  
         ...  
         Terrenth    1  
         Sparky      1  
         Spencer     1  
         Shadoe      1  
         Covach      1  
         Crumpet     1
```

Strider	1
Kobe	1
Tedders	1
Mitch	1
Alexander	1
Major	1
Lucky	1
Rizzo	1
Jennifur	1
Tobi	1
Mutt	1
Bronte	1
Holly	1
Ralphé	1
Murphy	1
Ozzie	1
Kulet	1
Danny	1
Rumpole	1
Hector	1
Kota	1
Kanu	1
Ricky	1
Pete	1

Name: name, Length: 932, dtype: int64

### 1.3.5 Issue #5

**Define:** The different dog stages should be made into a single column.

#### Code

```
In [122]: #adding all the columns into a single column
          archive['stages'] = archive.doggo + archive.floofer + archive.pupper + archive.puppo
```

```
In [125]: #checking to see twenty random samples of the new column
          archive.stages.sample(20)
```

```
Out[125]: 512      NoneNoneNoneNone
          1944      NoneNoneNoneNone
          1209      NoneNoneNoneNone
          2249      NoneNoneNoneNone
          138       NoneNoneNoneNone
          692       NoneNoneNoneNone
          399       NoneNoneNoneNone
          60        NoneNoneNoneNone
          1690      NoneNoneNoneNone
          1781      NoneNoneNoneNone
```

```

1149      NoneNoneNoneNone
1468      NoneNoneNoneNone
2230      NoneNoneNoneNone
501       doggoNoneNoneNone
1235      NoneNoneNoneNone
1434      NoneNoneNoneNone
2049      NoneNoneNoneNone
1248      NoneNoneNoneNone
1930      NoneNonepupperNone
207       NoneNoneNoneNone
Name: stages, dtype: object

```

In [126]: *#defining a function to extract the different stages*

```

def extract_stages(archive):
    if archive['stages'].count('None') == 2:
        return 'Multiple'
    else:
        if archive['stages'].count('doggo') == 1:
            return 'Doggo'
        elif archive['stages'].count('floofer') == 1:
            return 'Floofer'
        elif archive['stages'].count('pupper') == 1:
            return 'Pupper'
        elif archive['stages'].count('puppo') == 1:
            return 'Puppo'
        else:
            return 'None'

```

In [127]: archive['dog-stage'] = archive.apply(extract\_stages, axis=1)

In [143]: archive.info()

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 2356 entries, 0 to 2355
Data columns (total 7 columns):
tweet_id      2356 non-null object
timestamp     2356 non-null datetime64[ns]
text          2356 non-null object
rating_numerator    2356 non-null int64
rating_denominator  2356 non-null int64
name          2356 non-null object
dog-stage     2356 non-null object
dtypes: datetime64[ns](1), int64(2), object(4)
memory usage: 128.9+ KB

```

In [129]: *#Dropping the extra columns I no longer need; doggo, floofer, pupper, puppo and stages*  
archive.drop(['doggo', 'floofer', 'pupper', 'puppo', 'stages'], axis=1, inplace=True)

## Test

```
In [144]: archive.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 2356 entries, 0 to 2355
Data columns (total 7 columns):
tweet_id          2356 non-null object
timestamp         2356 non-null datetime64[ns]
text              2356 non-null object
rating_numerator  2356 non-null int64
rating_denominator 2356 non-null int64
name              2356 non-null object
dog-stage         2356 non-null object
dtypes: datetime64[ns](1), int64(2), object(4)
memory usage: 128.9+ KB
```

```
In [135]: archive['dog-stage'].value_counts()
```

```
Out[135]: None          1976
         Pupper         245
         Doggo          83
         Puppo          29
         Multiple       14
         Floofer         9
         Name: dog-stage, dtype: int64
```

### 1.3.6 Issue #6

**Define:** The links in the text column should be removed.

## Code

```
In [168]: archive['text'] = archive['text'].apply(lambda x: re.split('https://\/*', str(x))[0])
```

## Test

```
In [174]: for row in archive.text[:10]:
           print(row)
```

This is Phineas. He's a mystical boy. Only ever appears in the hole of a donut. 13/10

This is Tilly. She's just checking pup on you. Hopes you're doing ok. If not, she's available fo

This is Archie. He is a rare Norwegian Pouncing Corgo. Lives in the tall grass. You never know w

This is Darla. She commenced a snooze mid meal. 13/10 happens to the best of us

This is Franklin. He would like you to stop calling him "cute." He is a very fierce shark and sh

Here we have a majestic great white breaching off South Africa's coast. Absolutely h\*ckin breath

Meet Jax. He enjoys ice cream so much he gets nervous around it. 13/10 help Jax enjoy more thing

When you watch your owner call another dog a good boy but then they turn back to you and say you This is Zoey. She doesn't want to be one of the scary sharks. Just wants to be a snuggly pettable This is Cassie. She is a college pup. Studying international doggo communication and stick theorie

### 1.3.7 Issue #7:

**Define:** Convert the text in the p1, p2 and p3 columns of the images dataframe to lower case characters.

**Code:**

```
In [71]: images.p1 = images.p1.str.lower()
         images.p2 = images.p2.str.lower()
         images.p3 = images.p3.str.lower()
```

**Test**

```
In [72]: images.head()
```

```
Out[72]:
```

	tweet_id	jpg_url	\
0	666020888022790149	https://pbs.twimg.com/media/CT4udnOWwAAOaMy.jpg	
1	666029285002620928	https://pbs.twimg.com/media/CT42GRgUYAA5iDo.jpg	
2	666033412701032449	https://pbs.twimg.com/media/CT4521TWwAEvMyu.jpg	
3	666044226329800704	https://pbs.twimg.com/media/CT5Dr8HUEAA-lEu.jpg	
4	666049248165822465	https://pbs.twimg.com/media/CT5IQmsXIAAKY4A.jpg	

	img_num	p1	p1_conf	p1_dog	p2	\
0	1	welsh_springer_spaniel	0.465074	True	collie	
1	1	redbone	0.506826	True	miniature_pinscher	
2	1	german_shepherd	0.596461	True	malinois	
3	1	rhodesian_ridgeback	0.408143	True	redbone	
4	1	miniature_pinscher	0.560311	True	rottweiler	

	p2_conf	p2_dog	p3	p3_conf	p3_dog
0	0.156665	True	shetland_sheepdog	0.061428	True
1	0.074192	True	rhodesian_ridgeback	0.072010	True
2	0.138584	True	bloodhound	0.116197	True
3	0.360687	True	miniature_pinscher	0.222752	True
4	0.243682	True	doberman	0.154629	True

### 1.3.8 Tidiness

**Define:** Merge the dataframes into one.

## Code

```
In [176]: #merge the two tables
          archive = pd.merge(left=archive, right=tweets, how='left', on='tweet_id')
          archive = pd.merge(left=archive, right=images, how='left', on='tweet_id')
```

## Test

```
In [177]: archive.info()

<class 'pandas.core.frame.DataFrame'>
Int64Index: 2356 entries, 0 to 2355
Data columns (total 20 columns):
tweet_id          2356 non-null object
timestamp         2356 non-null datetime64[ns]
text              2356 non-null object
rating_numerator  2356 non-null int64
rating_denominator 2356 non-null int64
name              2356 non-null object
dog-stage         2356 non-null object
favorite_count    2327 non-null float64
retweet_count     2327 non-null float64
jpg_url           2075 non-null object
img_num           2075 non-null float64
p1                2075 non-null object
p1_conf           2075 non-null float64
p1_dog            2075 non-null object
p2                2075 non-null object
p2_conf           2075 non-null float64
p2_dog            2075 non-null object
p3                2075 non-null object
p3_conf           2075 non-null float64
p3_dog            2075 non-null object
dtypes: datetime64[ns](1), float64(6), int64(2), object(11)
memory usage: 386.5+ KB
```

```
In [179]: #drop rows with missing values from the merged dataframe
          archive.dropna(axis=0, inplace=True)
```

```
In [180]: archive.info()

<class 'pandas.core.frame.DataFrame'>
Int64Index: 2057 entries, 0 to 2355
Data columns (total 20 columns):
tweet_id          2057 non-null object
timestamp         2057 non-null datetime64[ns]
text              2057 non-null object
rating_numerator  2057 non-null int64
```



```

rating_denominator    2057 non-null int64
name                  2057 non-null object
dog-stage              2057 non-null object
favorite_count         2057 non-null float64
retweet_count          2057 non-null float64
jpg_url               2057 non-null object
img_num               2057 non-null float64
p1                    2057 non-null object
p1_conf               2057 non-null float64
p1_dog                2057 non-null object
p2                    2057 non-null object
p2_conf               2057 non-null float64
p2_dog                2057 non-null object
p3                    2057 non-null object
p3_conf               2057 non-null float64
p3_dog                2057 non-null object
dtypes: datetime64[ns](1), float64(6), int64(2), object(11)
memory usage: 337.5+ KB

```

## 1.4 Storing Data

Save gathered, assessed, and cleaned master dataset to a CSV file named "twitter\_archive\_master.csv".

```

In [ ]: archive.to_csv('twitter_archive_master.csv', index=False)

In [2]: master = pd.read_csv("twitter_archive_master.csv")
        master.info()

```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 2057 entries, 0 to 2056
Data columns (total 20 columns):
tweet_id              2057 non-null int64
timestamp             2057 non-null object
text                  2057 non-null object
rating_numerator      2057 non-null int64
rating_denominator    2057 non-null int64
name                  2057 non-null object
dog-stage             2057 non-null object
favorite_count         2057 non-null float64
retweet_count          2057 non-null float64
jpg_url               2057 non-null object
img_num               2057 non-null float64
p1                    2057 non-null object
p1_conf               2057 non-null float64
p1_dog                2057 non-null bool
p2                    2057 non-null object
p2_conf               2057 non-null float64

```

```

p2_dog          2057 non-null bool
p3              2057 non-null object
p3_conf         2057 non-null float64
p3_dog          2057 non-null bool
dtypes: bool(3), float64(6), int64(3), object(8)
memory usage: 279.3+ KB

```

- The favorite and retweet counts are floats upon conversion to csv, so these should be converted to int.
- The tweet\_id should also be re-converted to object datatype as well as timestamp to datatype datetime.
- The dog-stages column should be converted to category datatype.

```

In [4]: master['tweet_id'] = master['tweet_id'].astype(str)
        master['dog-stage'] = master['dog-stage'].astype("category")
        master[['rating_numerator', 'rating_denominator']] = master[['rating_numerator', 'rating_denominator']].astype(int)
        master[['retweet_count', 'favorite_count']] = master[['retweet_count', 'favorite_count']].astype(int)
        master.timestamp = pd.to_datetime(master.timestamp)

```

## 1.5 Analyzing and Visualizing Data

In this section, analyze and visualize your wrangled data. You must produce at least **three (3) insights and one (1) visualization**.

### 1.5.1 Insights

1. Is there any correlation between the retweet and favorite count?
2. What is the mean of the rating numerator?
3. What dog name is the most common?

### 1.5.2 Is there any correlation between the retweet and favorite count??

```

In [9]: master['favorite_count'].corr(master['retweet_count'])

```

```

Out[9]: 0.86104253667814268

```

The above shows a positive correlation between the retweet count and the favorite count.

### 1.5.3 What is the mean of the rating numerator?

```

In [103]: master.rating_numerator.describe()

```

```

Out[103]: count    2057.000000
          mean      12.265921
          std       40.857799
          min        0.000000

```

```


25%      10.000000
50%      11.000000
75%      12.000000
max      1776.000000
Name: rating_numerator, dtype: float64

```

As seen here, the mean for the rating numerator is 12.

#### 1.5.4 What dog name is the most common?

```
In [6]: master.name.value_counts()[1:10].plot(kind='bar');
```



As seen above, for this dataset, the most common dog names are Tucker, Cooper, Oliver, Penny and Charlie.

#### 1.5.5 Visualization

```

In [26]: import seaborn as sns
plt.figure(figsize=(8,8))
sns.lmplot(x='retweet_count', y='favorite_count', data=master,fit_reg=True)
plt.xlabel("Retweet Count")
plt.ylabel("Favorite Count")
plt.title("Correlation between Retweet and Favorite Count");

```

```
<matplotlib.figure.Figure at 0x7f6062e05c88>
```

